

```
1 #include <fcntl.h>
2 #include <netinet/in.h>
3 #include <unistd.h>
4 #include <signal.h>
5 #include <sys/ioctl.h>
6 #include <sys/socket.h>
7 #include <sys/time.h>
8
9 #include <atomic>
10 #include <cerrno>
11 #include <cstdio>
12 #include <cstdlib>
13 #include <vector>
14
15 using namespace std;
16
17 #define MSG_NOSIGNAL 0
18
19 // Driver stuff
20 #define CSGAMES_MAX_STARTUPS 8
21 #define CSGAMES_MAX_STARTUP_NAME_SIZE 100 // including NUL terminator
22 #define CSGAMES_STARTUP_PREDICTED_YEAR_COUNT 16
23
24 #define SIGNGEL SIGUSR1
25
26 struct startup {
27     char name[CSGAMES_MAX_STARTUP_NAME_SIZE];
28     unsigned capital_in_dollars;
29 };
30
31 struct startup_net_worth {
32     startup startup;
33     long worth_by_year[CSGAMES_STARTUP_PREDICTED_YEAR_COUNT];
34 };
35
36 enum csgames_ioctl_request_code {
37     CSGAMES_NEW_STARTUP, // (const struct startup*)
38     CSGAMES_SET_STARTUP, // (const char*)
39     CSGAMES_GET_PROCESSING_END, // (timeval*)
40     CSGAMES_GET_ANGELS_PICK, // (char [CSGAMES_MAX_STARTUP_NAME_SIZE])
41 };
42
43 // Socket stuff
44 const sockaddr_in bind_addr = {
45     .sin_family = AF_INET,
46     .sin_port = htons(4321),
47     .sin_addr = { htonl(INADDR_LOOPBACK) },
48 };
49
```

```
50 // Write all numeric values in network byte order!
51 enum csgames_protocol_command_code {
52     BEGIN_STARTUP,          // (uint32_t, char [])
53     BUSINESS_PLAN_DATA,     // (const char [])
54
55     ACKNOWLEDGED,           // (void)
56     ANGEL_INVESTOR,         // (void)
57     PREDICTION_RESULT,      // (long [CSGAMES_STARTUP_PREDICTED_YEAR_COUNT])
58     ERROR,                  // (const char [])
59 };
60
61 // Program
62 struct startup_entry {
63     startup startup;
64     timeval ready_at;
65     pid_t pid;
66 };
67
68 startup_entry all_startups[CSGAMES_MAX_STARTUPS];
69 atomic_int startup_count;
70 int ss;
71 int lock;
72
73 struct command_header {
74     unsigned short command;
75     unsigned short size;
76 };
77
78 struct begin_startup_command {
79     command_header header;
80     unsigned capital;
81     char name[100];
82 };
83
84 struct business_plan_data_command {
85     command_header header;
86     uint8_t data[1];
87 };
88
89 struct prediction_result {
90     command_header header;
91     long values[CSGAMES_STARTUP_PREDICTED_YEAR_COUNT];
92 };
93
94 union command_buffer {
95     command_header header;
96     begin_startup_command begin_startup;
97     business_plan_data_command business_plan_data;
98     prediction_result prediction;
```

```
99     uint8_t bytes[0xffff + sizeof(command_header)];
100 };
101
102 struct lock_file {
103     int fd;
104     lock_file(int fd) : fd(fd) { flock(fd, LOCK_EX); }
105     ~lock_file() { flock(fd, LOCK_UN); }
106 };
107
108 int read_exactly(int sock, void* buffer, ssize_t count) {
109     uint8_t* cBuffer = reinterpret_cast<uint8_t*>(buffer);
110     ssize_t readTotal = 0;
111     while (readTotal != count) {
112         ssize_t result = recv(sock, cBuffer + readTotal, count - readTotal, MSG_NOSIGNAL);
113         if (result < 0 && errno != EINTR) {
114             return errno;
115         }
116         readTotal += result;
117     }
118     return 0;
119 }
120
121 int send_exactly(int sock, const void* buffer, ssize_t count) {
122     const uint8_t* cBuffer = reinterpret_cast<const uint8_t*>(buffer);
123     ssize_t sentTotal = 0;
124     while (sentTotal != count) {
125         ssize_t result = send(sock, cBuffer + sentTotal, count - sentTotal, MSG_NOSIGNAL);
126         if (result < 0 && errno != EINTR) {
127             return errno;
128         }
129         sentTotal += result;
130     }
131     return 0;
132 }
133
134 bool before(const timeval& a, const timeval& b) {
135     return a.tv_sec < b.tv_sec || (a.tv_sec == b.tv_sec && a.tv_usec < b.tv_usec);
136 }
137
138 int send_error(int client, const char* fmt, ...) {
139     char* message;
140
141     va_list ap;
142     va_start(ap, fmt);
143     vasprintf(&message, fmt, ap);
144     va_end(ap);
145
146     size_t len = strlen(message);
147     command_header header = {
```

```
148     .command = htons(ERROR),
149     .size = htons(sizeof(command_header) + len)
150 };
151
152 int result = send_exactly(client, &header, sizeof(header));
153 if (result == 0) {
154     result = send_exactly(client, &message, len);
155 }
156 free(message);
157 return result;
158 }
159
160 void serve_client(int client, startup_entry* entry) {
161     vector<uint8_t> bytes;
162
163     // read input
164     while (true) {
165         command_buffer buffer;
166         int result = read_exactly(client, buffer.bytes, sizeof buffer.header);
167         if (result != 0) {
168             bytes.clear();
169             break;
170         }
171
172         buffer.header.command = ntohs(buffer.header.command);
173         buffer.header.size = ntohs(buffer.header.size);
174         if (buffer.header.size != 0) {
175             result = read_exactly(client, buffer.business_plan_data.data, buffer.header.size);
176             if (result != 0) {
177                 bytes.clear();
178                 break;
179             }
180         }
181
182         if (buffer.header.command == BUSINESS_PLAN_DATA) {
183             command_header ack = { ACKNOWLEDGED, sizeof (command_header) };
184             result = send_exactly(client, &ack, sizeof ack);
185             if (result != 0) {
186                 bytes.clear();
187                 break;
188             }
189             if (buffer.header.size == 0) {
190                 bytes.clear();
191                 break;
192             }
193             else {
194                 bytes.insert(bytes.end(), buffer.business_plan_data.data, buffer.bytes.end());
195             }
196         }
197     }
198 }
```

```
197     else {
198         result = send_error(client, "expected BUSINESS_PLAN_DATA command (%i), g
199     }
200 }
201
202 // drain buffer
203 timeval now = {};
204 size_t index = 0;
205 while (index != bytes.size()) {
206     gettimeofday(&now, nullptr);
207     if (before(now, entry->ready_at)) {
208         useconds_t sleep_time = entry->ready_at.tv_sec - now.tv_sec;
209         sleep_time *= 1000000;
210         sleep_time += entry->ready_at.tv_usec - now.tv_usec;
211         usleep(sleep_time);
212         continue;
213     }
214
215     lock_file l(lock);
216     int result = ioctl(ss, CSGAMES_SET_STARTUP, entry->startup.name);
217     if (result < 0) {
218         if (errno == EINTR) {
219             continue;
220         } else {
221             break;
222         }
223     }
224     result = write(ss, &bytes[index], bytes.size() - index);
225     if (result < 0) {
226         if (errno != EAGAIN) {
227             break;
228         }
229     } else {
230         index += result;
231     }
232
233     result = ioctl(ss, CSGAMES_GET_PROCESSING_END, &entry->ready_at);
234     if (result < 0) {
235         break;
236     }
237 }
238
239 // read and send back result
240 startup_net_worth output;
241 {
242     lock_file l(lock);
243     int result = ioctl(ss, CSGAMES_SET_STARTUP, entry->startup.name);
244     if (result < 0) {
245         return;
```

```
246     }
247     result = read(ss, &output, sizeof output);
248     if (result < 0) {
249         return;
250     }
251 }
252
253 prediction_result result;
254 result.header.command = htons(PREDICTION_RESULT);
255 result.header.size = htons(sizeof result);
256 for (int i = 0; i < CSGAMES_STARTUP_PREDICTED_YEAR_COUNT; ++i) {
257     result.values[i] = htonl(output.worth_by_year[i]);
258 }
259 send_exactly(client, &result, sizeof result);
260 }
261
262 void handle_client(int client) {
263     command_buffer buffer = {};
264     int result;
265
266     // THING TO TEST:
267     // How do servers react when clients block here? (This implementation does
268     // not do the right thing.)
269     do
270     {
271         result = read_exactly(client, buffer.bytes, sizeof buffer.header);
272         if (result != 0) {
273             fprintf(stderr, "error %i reading command header!\n", result);
274             return;
275         }
276
277         buffer.header.command = htons(buffer.header.command);
278         buffer.header.size = htons(buffer.header.size);
279         if (buffer.header.command != BEGIN_STARTUP) {
280             send_error(client, "expected BEGIN_STARTUP!");
281             continue;
282         }
283         if (buffer.header.size > sizeof(begin_startup_command)) {
284             // THING TO TEST:
285             // How do servers react to a buffer that is too large?
286             send_error(client, "BEGIN_STARTUP command is too large!");
287             continue;
288         }
289
290         result = read_exactly(client, buffer.bytes + sizeof buffer.header, buffer.header.size);
291         if (result != 0) {
292             fprintf(stderr, "error %i reading command body!\n", result);
293             return;
294         }
295     }
```

```
295     buffer.begin_startup.capital = htonl(buffer.begin_startup.capital);
296     buffer.bytes[buffer.header.size] = 0; // nullptr-terminate buffer
297     break;
298 } while (true);
299
300 pid_t pid = -1;
301 sigset_t set;
302 sigprocmask(0, nullptr, &set);
303 sigset_t oldMask = set;
304
305 sigaddset(&set, SIGCHLD);
306 sigaddset(&set, SINGEL);
307 sigprocmask(SIG_SETMASK, &set, nullptr);
308
309 int index = startup_count;
310 memset(&all_startups[index], 0, sizeof all_startups[index]);
311 all_startups[index].startup.capital_in_dollars = buffer.begin_startup.capital;
312
313 size_t nameSize = buffer.begin_startup.header.size
314     - sizeof(buffer.begin_startup.header)
315     - sizeof(buffer.begin_startup.capital);
316 strncpy(all_startups[index].startup.name, buffer.begin_startup.name, sizeof(all
317
318 result = ioctl(ss, CSGAMES_NEW_STARTUP, &all_startups[index].startup);
319 if (result == 0)
320 {
321     pid = fork();
322     if (pid == -1) {
323         perror("fork");
324     }
325     else if (pid != 0) {
326         all_startups[index].pid = pid;
327         ++startup_count;
328     }
329 }
330 sigprocmask(SIG_SETMASK, &oldMask, nullptr);
331
332 if (pid == 0) {
333     serve_client(client, &all_startups[index]);
334 }
335 }
336
337 void sigchld(int, siginfo_t* info, void*) {
338     for (int i = 0; i < startup_count; ++i) {
339         if (info->si_pid == all_startups[i].pid) {
340             memmove(
341                 &all_startups[i],
342                 &all_startups[i+1],
343                 sizeof(all_startups[0]) * (CSGAMES_MAX_STARTUPS - i - 1));
```

```
344         --startup_count;
345         break;
346     }
347 }
348 }
349
350 void signgel(int) {
351     char startup_name[CSGAMES_MAX_STARTUP_NAME_SIZE];
352     int err = ioctl(ss, CSGAMES_GET_ANGELS_PICK, &startup_name);
353     if (err != 0) {
354         static char message[] = "couldn't get angel's pick!\n";
355         write(STDERR_FILENO, message, sizeof(message)-1);
356         exit(1);
357     }
358
359     for (int i = 0; i < startup_count; ++i) {
360         if (strncmp(startup_name, all_startups[i].startup.name, sizeof(startup_name))
361             kill(all_startups[i].pid, SIGQUIT);
362         return;
363     }
364 }
365 static char message[] = "couldn't find startup!\n";
366 write(STDERR_FILENO, message, sizeof(message)-1);
367 }
368
369 int main() {
370     int result;
371
372     #pragma mark - Signal Handlers
373     fprintf(stderr, "Setting up signal handlers... ");
374     struct sigaction chldAction;
375     result = sigaction(SIGCHLD, nullptr, &chldAction);
376     if (result != 0) {
377         perror("sigaction(SIGCHLD)");
378         return 1;
379     }
380     chldAction.sa_sigaction = &sigchld;
381     sigaddset(&chldAction.sa_mask, SIGNGEL);
382     sigaddset(&chldAction.sa_mask, SIGCHLD);
383     result = sigaction(SIGCHLD, &chldAction, nullptr);
384     if (result != 0) {
385         perror("sigaction(SIGCHLD)");
386         return 1;
387     }
388
389     struct sigaction ngelAction;
390     result = sigaction(SIGNGEL, nullptr, &ngelAction);
391     if (result != 0) {
392         perror("sigaction(SIGNGEL)");
```



```
393     return 1;
394 }
395 ngelAction.sa_handler = &signgel;
396 sigaddset(&ngelAction.sa_mask, SIGNGEL);
397 sigaddset(&ngelAction.sa_mask, SIGCHLD);
398 result = sigaction(SIGNGEL, &ngelAction, nullptr);
399 if (result != 0) {
400     perror("sigaction(SIGNGEL)");
401     return 1;
402 }
403 fprintf(stderr, "done\n");
404
405 #pragma mark - Preparing lock
406 char lockFileName[] = "csgames_startup_simulator.XXXXXX";
407 lock = mkstemp(lockFileName);
408
409 #pragma mark - Opening sockets
410 int acceptor = socket(PF_INET, SOCK_STREAM, IPPROTO_TCP);
411 if (acceptor < 0) {
412     perror("socket");
413     return 1;
414 }
415
416 int yes = 1;
417 result = setsockopt(acceptor, SOL_SOCKET, SO_REUSEADDR, &yes, sizeof(yes));
418 if (result != 0) {
419     perror("setsockopt");
420     return 1;
421 }
422
423 result = bind(acceptor, (const struct sockaddr*)&bind_addr, sizeof(bind_addr));
424 if (result != 0) {
425     perror("bind");
426     return 1;
427 }
428
429 result = listen(acceptor, 24);
430 if (result != 0) {
431     perror("listen");
432     return 1;
433 }
434
435 while (1) {
436     // THING TO TEST:
437     // how do servers handle more than 8 clients?
438     while (startup_count == 8) {
439         pause();
440     }
441 }
```

```
442     struct sockaddr_in client_address;
443     socklen_t client_address_length = sizeof(client_address);
444     int client = accept(acceptor, (struct sockaddr*)&client_address, &client_ad
445     if (client < 0) {
446         perror("accept");
447         continue;
448     }
449
450     // THINGS TO TEST:
451     // How does a client handle being stuck waiting?
452     handle_client(client);
453     close(client);
454 }
455 }
456
```